

```
#!/pip install pyspark
```

```
Collecting pyspark
  Downloading pyspark-3.5.1.tar.gz (317.0 MB)
    317.0/317.0 MB 3.9 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.10/dist-packages (from pyspark) (0.10.9.7)
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.5.1-py2.py3-none-any.whl size=317488493 sha256=a8fad28035e732e4cd3556eaa6d4b
  Stored in directory: /root/.cache/pip/wheels/80/1d/60/2c256ed38dddce2fdd93be545214a63e02fbd8d74fb0b7f3a6
Successfully built pyspark
Installing collected packages: pyspark
Successfully installed pyspark-3.5.1
```

✓ Mount Shared Drive

Mount a google shared drive in order to develop ETL for model development

```
import pandas as pd
from google.colab import drive
drive.mount('/content/gdrive')

Mounted at /content/gdrive
```

✓ Create Spark Session

Create a spark session called spotify which will handle Spark functionalities.

```
from pyspark.sql import SparkSession

# Create Spark Environment
ss = SparkSession.builder.appName("spotify").getOrCreate()
```

✓ Sample Playlist Data

The dataset filled with a million playlists consists of 33.54 GB of json data files. On a cluster of servers, it should not be an issue, however, we are limited to a pro version of Google Colab which cannot handle the large dataset. We, therefore, need to resample the dataset to obtain a smaller working size. The sample size chosen was 5 files consisting of the portion of the playlist within the dataset. This, however, is a more manageable data size for our computational power.

```

import os
import random
import re
import shutil

#===== Sample Data =====
def random_sample_from_folder(folder_path, sample_size=10):
    """
    Select a random sample of files from the specified folder.

    :param folder_path: Path to the folder from which to sample files.
    :param sample_size: Number of files to sample.
    :return: A list of paths to the randomly sampled files.
    """
    # List all files in the folder
    all_files = [os.path.join(folder_path, f) for f in os.listdir(folder_path) if os.path.isfile(os.path.join(folder_path, f))]

    # Determine the sample size (can't be larger than the number of files)
    actual_sample_size = min(sample_size, len(all_files))

    # Randomly sample files
    sampled_files = random.sample(all_files, actual_sample_size)

    return sampled_files

# Create Sample
folder_path = "/content/gdrive/Shareddrives/737_Project/spotify_million_playlist_dataset/data/"
sample_size = 5 # Adjust as needed
sampled_files = random_sample_from_folder(folder_path, sample_size)

#===== Abstract File Name =====
# Regular expression pattern to match the file name at the end of the path
pattern = r'(/[^\s]+\.[json])$'

# Initialize a list to hold the matched file names
matched_file_names = []

# Iterate over the file paths and match the pattern
for file_path in sampled_files:
    match = re.search(pattern, file_path)
    if match:
        matched_file_names.append(match.group(1)) # Append the matched file name

# Print all matched file names
# for file_name in matched_file_names:
#     print(file_name)

#print(len(sampled_files))

#===== Move Data To Sampled Path =====
source_folder = '/content/gdrive/Shareddrives/737_Project/spotify_million_playlist_dataset/data/'
destination_folder = '/content/gdrive/Shareddrives/737_Project/spotify_million_playlist_dataset/sampled_data'

# Create New File
if not os.path.exists(destination_folder):
    os.makedirs(destination_folder)

# Move each file to the destination folder
for file in matched_file_names:
    # Construct the full file path
    source_file_path = os.path.join(source_folder, file)
    destination_file_path = os.path.join(destination_folder, file)

    # Move the file
    shutil.move(source_file_path, destination_file_path)

5

```

✓ Load sampled data in to Pyspark

This loads our data files within a sampled folder and reads it to a spark data frame. It can handle json file as seen in the "View of sampled data files by track_uri" within the next block

```
# Folder Path of Playlist
folder_path = "/content/gdrive/SharedDrives/737_Project/spotify_million_playlist_dataset/sampled_data/*"

df_playlist = ss.read.option("multiline", "true").json(folder_path)
#df = ss.read.option("mergeSchema", "true").json(folder_path)
```

✓ View of sampled data files by track_uri

Here we can see that the files have been properly loaded as we can witness pid and track_uri. Pid is the playlist id and the track_uri is the track identifier. This is for future implementations to build out song recommendations for playlists. The current implementation only has recommendations based on a single song.

```
from pyspark.sql.functions import explode, col
from pyspark.sql.functions import lit
tracks_df = df_playlist.select(explode("playlists").alias("playlist")) \
    .select("playlist.pid", explode("playlist.tracks").alias("track")) \
    .select("pid", "track.track_uri")

#tracks_df = tracks_df.withColumn("rating", lit(1))
tracks_df.show()
```

```
+-----+-----+
|  pid|          track_uri|
+-----+-----+
|150000|spotify:track:5CB...|
|150000|spotify:track:4Sn...|
|150000|spotify:track:50Y...|
|150000|spotify:track:636...|
|150000|spotify:track:2ay...|
|150000|spotify:track:5IM...|
|150000|spotify:track:5C9...|
|150000|spotify:track:4tD...|
|150000|spotify:track:15W...|
|150000|spotify:track:2kD...|
|150000|spotify:track:3yt...|
|150000|spotify:track:50K...|
|150000|spotify:track:0TY...|
|150000|spotify:track:0fy...|
|150000|spotify:track:2bh...|
|150001|spotify:track:76V...|
|150001|spotify:track:6eF...|
|150001|spotify:track:6Jn...|
|150001|spotify:track:0CK...|
|150001|spotify:track:7dC...|
+-----+-----+
only showing top 20 rows
```

✓ Recommend Users who like a track to an Artist

✓ Read Music Data

Load the data into spark and clean out columns that won't be need in model development.

```
from pyspark.sql.types import StructType, StructField, IntegerType, StringType, FloatType, BooleanType

# Load dataset
spotify_song_df = ss.read.csv("/content/gdrive/SharedDrives/737_Project/spotify_million_playlist_dataset/spotify_data.csv", head=1, inferSchema=True)
col_drop = ['_c0', 'key', 'duration_ms', 'time_signature'] #['Unnamed: 0', 'artist_name', 'track_name', 'key', 'duration_ms', 'time_signature']
spotify_song_df = spotify_song_df.drop(*col_drop)
spotify_song_df.printSchema()

root
 |-- track_id: string (nullable = true)
 |-- artists: string (nullable = true)
 |-- album_name: string (nullable = true)
 |-- track_name: string (nullable = true)
 |-- popularity: string (nullable = true)
```

```

|-- explicit: string (nullable = true)
|-- danceability: string (nullable = true)
|-- energy: string (nullable = true)
|-- loudness: string (nullable = true)
|-- mode: string (nullable = true)
|-- speechiness: string (nullable = true)
|-- acousticness: string (nullable = true)
|-- instrumentalness: double (nullable = true)
|-- liveness: string (nullable = true)
|-- valence: string (nullable = true)
|-- tempo: double (nullable = true)
|-- track_genre: string (nullable = true)

```

```
spotify_song_df.show(truncate = False)
```

id	artists	album_name	track_name
1RyPMVoIQDJUgSV	Gen Hoshino	Comedy	Comedy
1i3p13qLCt0Ki3A	Ben Woodward	Ghost (Acoustic)	Ghost - Acoustic
s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN	To Begin Again	To Begin Again
G4xtTiEg7opyCyx	Kina Grannis	Crazy Rich Asians (Original Motion Picture Soundtrack)	Can't Help Fallin
imiIP26QG5WcN2K	Chord Overstreet	Hold On	Hold On
KtVTNFFiBU9I7dc	Tyrone Wells	Days I Will Remember	Days I Will Remer
lmXdkIAM7WUoEb7N	A Great Big World;Christina Aguilera	Is There Anybody Out There?	Say Something
imMH3G43AXT1y7pA	Jason Mraz	We Sing. We Dance. We Steal Things.	I'm Yours
nAGrvD03AWnz3Q8	Jason Mraz;Colbie Caillat	We Sing. We Dance. We Steal Things.	Lucky
Lp2AzqokyEdwEw2	Ross Copperman	Hunger	Hunger
IkRvGxdhdGdAH7EJ	Zack Tabudlo	Episode	Give Me Your For
lBqJiVL5IAE9jRyl	Jason Mraz	Love Is a Four Letter Word	I Won't Give Up
l35d7gQfeNteBwp	Dan Berk	Solo	Solo
1rTkEHDjp95F200	Anna Hamilton	Bad Liar	Bad Liar
N82ZRhz9jqzgrb3	Chord Overstreet;Deepend	Hold On (Remix)	Hold On - Remix
IK9QxnGjdXb55NiG	Landon Pigg	The Boy Who Never	Falling in Love ;
fjixSUld14qUezm	Andrew Foy;Renee Foy	ily (i love you baby)	ily (i love you l
icoNyat1secaH00D	Andrew Foy;Renee Foy	At My Worst	At My Worst
uEC3ruGJg45MMN6	Jason Mraz;Colbie Caillat	We Sing. We Dance. We Steal Things.	Lucky
FJ4Q4Id4EjtbXLC	Boyce Avenue;Bea Miller	Cover Sessions, Vol. 4	Photograph

```
wing top 20 rows
```

✓ ALS Model For Artist Recommendation Based on Popularity

The ALS model develops a recommendation system for a user who likes a song to explore other artists that are of the same popularity. The model takes into consideration the track_id, the artist_id, and popularity column to offer the top 10 recommended artist for a user who likes a song.

```

from pyspark.ml.feature import StringIndexer, IndexToString, OneHotEncoder, VectorAssembler
from pyspark.ml import Pipeline
from pyspark.ml.recommendation import ALS
from pyspark.ml.evaluation import RegressionEvaluator

# Convert columns to data types
df = spotify_song_df \
    .withColumn('popularity', spotify_song_df['popularity'].cast('int')) \
    .withColumn('explicit', spotify_song_df['explicit'].cast('boolean')) \
    .withColumn('danceability', spotify_song_df['danceability'].cast('double')) \
    .withColumn('energy', spotify_song_df['energy'].cast('double')) \
    .withColumn('loudness', spotify_song_df['loudness'].cast('double')) \
    .withColumn('mode', spotify_song_df['mode'].cast('int')) \
    .withColumn('speechiness', spotify_song_df['speechiness'].cast('double')) \
    .withColumn('acousticness', spotify_song_df['acousticness'].cast('double')) \
    .withColumn('instrumentalness', spotify_song_df['instrumentalness'].cast('double')) \
    .withColumn('liveness', spotify_song_df['liveness'].cast('double')) \
    .withColumn('valence', spotify_song_df['valence'].cast('double')) \
    .withColumn('tempo', spotify_song_df['tempo'].cast('double'))

# Drop rows with NaN values in any column
spotify_song_df = spotify_song_df.dropna(how='any')

# StringIndexer to convert string fields to indices which will be used by ALS
indexers = [StringIndexer(inputCol=column, outputCol=column+"_index").fit(df) for column in list(set(df.columns)-set(['popularity', 'track_id', 'artists', 'album_name', 'track_name', 'track_genre']))]

# Index the genre column
genre_indexer = StringIndexer(inputCol='track_genre', outputCol='track_genre_index').fit(df)

# One-hot encode the indexed genre column
genre_encoder = OneHotEncoder(inputCol='track_genre_index', outputCol='track_genre_vec')

assembler=VectorAssembler(inputCols=['track_id', 'artists', 'album_name', 'track_name', 'popularity', \
    'explicit', 'danceability', 'energy', 'loudness', 'mode', 'speechiness', \
    'acousticness', 'instrumentalness', 'liveness', 'valence', 'track_genre_vec'], outputCol='features')

# Define the ALS algorithm
als = ALS(
    userCol="track_id_index",
    itemCol="artists_index",
    ratingCol="popularity",
    nonnegative=True,
    implicitPrefs=False,
    coldStartStrategy="drop"
)

# Now, add these to your pipeline
pipeline_stages = indexers + [genre_indexer, genre_encoder, als]
pipeline = Pipeline(stages=pipeline_stages)

# Split the data into training and test sets
(training_data, test_data) = df.randomSplit([0.8, 0.2])

# Fit the model
model = pipeline.fit(training_data)

# Make predictions
predictions = model.transform(test_data)

# Evaluate the model by computing the RMSE on the test data
evaluator = RegressionEvaluator(metricName="rmse", labelCol="popularity", predictionCol="prediction")

rmse = evaluator.evaluate(predictions)

print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)

Root Mean Squared Error (RMSE) on test data = 2.04544

This portion recommends a user 10 artists based on songs they like.

# Get 10 artist to recommendation for a track
user_recs = model.stages[-1].recommendForAllUsers(10)

user_recs.show(truncate=False)

```

track_id_index	recommendations
26	[[{29461, 136.20064}, {17011, 114.20156}, {18244, 113.6354}, {24863, 112.77322}, {12008, 106.92717}, {7920, 1
27	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
28	[[{9134, 100.34088}, {10074, 98.84116}, {27974, 97.190216}, {30972, 95.44558}, {27411, 93.40673}, {23027, 90.
31	[[{5053, 108.15918}, {3788, 104.418076}, {25341, 101.096504}, {9634, 100.35293}, {4003, 98.2284}, {5735, 96.4
34	[[{5053, 117.77333}, {3788, 113.69968}, {25341, 110.08286}, {9634, 109.273186}, {4003, 106.959816}, {5735, 10
44	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
53	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
65	[[{23088, 14.280181}, {5052, 13.922445}, {30797, 13.549765}, {23716, 13.401429}, {29804, 13.358647}, {29461,
76	[[{13329, 78.078186}, {27316, 77.44498}, {10676, 75.35704}, {3865, 72.996315}, {6580, 71.56297}, {5771, 71.50
78	[[{16410, 74.35415}, {24864, 66.8762}, {8580, 66.67165}, {7307, 64.094955}, {17388, 62.098248}, {30970, 61.52
81	[[{29461, 164.05984}, {17011, 137.56097}, {18244, 136.879}, {24863, 135.84045}, {12008, 128.79865}, {7920, 12
85	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
101	[[{6234, 94.50512}, {16366, 87.20385}, {16001, 84.42688}, {19243, 83.72504}, {16280, 83.58065}, {1561, 83.436
103	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
108	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
115	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
126	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
133	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0
137	[[{29461, 164.05984}, {17011, 137.56097}, {18244, 136.879}, {24863, 135.84045}, {12008, 128.79865}, {7920, 12
148	[[{99, 0.0}, {98, 0.0}, {97, 0.0}, {96, 0.0}, {95, 0.0}, {94, 0.0}, {93, 0.0}, {92, 0.0}, {91, 0.0}, {90, 0.0

only showing top 20 rows

✓ Song Recomender

The song recommender recommends songs to a user based on a song that they like. The model is built using cosine similarity which takes into account variables on aspects of the song to recommend similar ones. The model takes a song name and recommends 10 other songs that the user might like.

```
from pyspark.ml.feature import VectorAssembler, StringIndexer, OneHotEncoder, Normalizer
from pyspark.ml import Pipeline
```

```
# Drop NaNs
```

```
spotify_song_df = spotify_song_df.dropna(how='any')
```

```
# Convert columns to data types
```

```
spotify_song_df = spotify_song_df \
    .withColumn('popularity', spotify_song_df['popularity'].cast('int')) \
    .withColumn('explicit', spotify_song_df['explicit'].cast('boolean')) \
    .withColumn('danceability', spotify_song_df['danceability'].cast('double')) \
    .withColumn('energy', spotify_song_df['energy'].cast('double')) \
    .withColumn('loudness', spotify_song_df['loudness'].cast('double')) \
    .withColumn('mode', spotify_song_df['mode'].cast('int')) \
    .withColumn('speechiness', spotify_song_df['speechiness'].cast('double')) \
    .withColumn('acousticness', spotify_song_df['acousticness'].cast('double')) \
    .withColumn('instrumentalness', spotify_song_df['instrumentalness'].cast('double')) \
    .withColumn('liveness', spotify_song_df['liveness'].cast('double')) \
    .withColumn('valence', spotify_song_df['valence'].cast('double')) \
    .withColumn('tempo', spotify_song_df['tempo'].cast('double'))
```

```
# Example of handling categorical variables
```

```
features = ['popularity', 'danceability', 'energy', 'loudness', 'mode', 'speechiness',
            'acousticness', 'instrumentalness', 'liveness', 'valence', 'tempo']
```

```
assembler_audio = VectorAssembler(inputCols=features, outputCol="audio_vec")
```

```
normalizer = Normalizer(inputCol="audio_vec", outputCol="normFeatures")
```

```
stringIndexer = StringIndexer(inputCol="track_genre", outputCol="genreIndexed")
```

```
encoder = OneHotEncoder(inputCols=["genreIndexed"], outputCols=["genreVec"])
```

```
# Assuming other features are numerical and already in the desired format
```

```
assembler = VectorAssembler(inputCols=["normFeatures", "genreVec"], outputCol="features")
```

```
pipeline = Pipeline(stages=[assembler_audio, normalizer, stringIndexer, encoder, assembler])
```

```
pipelineModel = pipeline.fit(spotify_song_df)
```

```
df_transformed = pipelineModel.transform(spotify_song_df)
```

```
df_transformed.show(truncate=False)
```

	genreIndexed	genreVec	features
5245735512779771,0.7679808798280549]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,5,6,7,8,9,10,108], [0.637676
5931002454,0.002764393570075794,0.8022849938262292]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.5694
53013726232395,0.7970420312564265]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.595181
339925472E-4,7.295997626585927E-4,0.9272549710879207]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.3622
52746553164,0.8236709262653021]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.563081
5829743383356447,0.8579789147243977]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.507695
368905538E-4,4.789025993661266E-4,0.8844611091352135]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.4632
34161074576565257,0.88224131752569]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.467536
30446231128999991,0.8677027669559167]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.493588
32554E-4,0.0020206501186201442,0.8134044576990345]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.5773
5761461562928,0.8018161090280541]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.593908
55783673381E-4,0.8862050960264252]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.458361
3.920255946618441]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,5,6,8,9,10,108], [0.38518689
3019510897373836735,0.813137652609235]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.578792
540146862E-4,0.002919085477163931,0.9051729550428433]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.4223
416635567,0.002334490699213771,0.8186117238835449]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.5689
55818911745,0.0033146713992235695,0.8819427857731233]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.4483
35636809091615388,0.8497142611454308]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,7,8,9,10,108], [0.499815
549368256869812,0.8846311177872647]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.462417
25287626764338403,0.8477323474233768]	97.0	(113, [97], [1.0])	(124, [0,1,2,3,4,5,6,8,9,10,108], [0.526171

Data cleaning process to convert a sparse vector to an array and then to a list. Needed to obtain a cosine similarity score as we will need to take the dot product of the vectors.

```
from pyspark.sql.types import ArrayType, DoubleType
from pyspark.sql.functions import udf
```

```
# UDF to convert a sparse vector to a dense vector
```

```
def sparse_to_dense(sparse_vector):
    return sparse_vector.toArray().toList()
```

```
# Register the UDF with the appropriate return type
```

```
sparse_to_array_udf = udf(sparse_to_dense, ArrayType(DoubleType()))
```

```
# Apply the UDF to the normalized feature column
```

```
df_transformed = df_transformed.withColumn("features", sparse_to_array_udf("features"))
```

```
df_transformed.show(truncate=False)
```

	genreIndexed	genreVec	features
5735512779771,0.7679808798280549]	97.0	(113, [97], [1.0])	[0.6376764929131795, 0.005905059030264512,
1002454,0.002764393570075794,0.8022849938262292]	97.0	(113, [97], [1.0])	[0.5694443683676729, 0.004348484267534956,
13726232395,0.7970420312564265]	97.0	(113, [97], [1.0])	[0.5951815199603877, 0.004573500100748242,
925472E-4,7.295997626585927E-4,0.9272549710879207]	97.0	(113, [97], [1.0])	[0.3622488332081125, 0.00135715760046983, 3
46553164,0.8236709262653021]	97.0	(113, [97], [1.0])	[0.5630811090859847, 0.00424370884652608, 0
9743383356447,0.8579789147243977]	97.0	(113, [97], [1.0])	[0.5076953697217327, 0.0060223174891129655,
905538E-4,4.789025993661266E-4,0.8844611091352135]	97.0	(113, [97], [1.0])	[0.46325218762213555, 0.002547887031921745,
61074576565257,0.88224131752569]	97.0	(113, [97], [1.0])	[0.46753646927699516, 0.0041084767237715945
46231128999991,0.8677027669559167]	97.0	(113, [97], [1.0])	[0.49358899171897364, 0.004168825943572412,
54E-4,0.0020206501186201442,0.8134044576990345]	97.0	(113, [97], [1.0])	[0.5773286053200412, 0.004556772206276039,
1461562928,0.8018161090280541]	97.0	(113, [97], [1.0])	[0.5939081334074972, 0.005032167562790551,
83673381E-4,0.8862050960264252]	97.0	(113, [97], [1.0])	[0.45836133026867854, 0.00320852931188075,
20255946618441]	97.0	(113, [97], [1.0])	[0.38518689911102383, 0.0036222383397171275
9510897373836735,0.813137652609235]	97.0	(113, [97], [1.0])	[0.5787921708985061, 0.006450732098239801,
146862E-4,0.002919085477163931,0.9051729550428433]	97.0	(113, [97], [1.0])	[0.4223999656878039, 0.005694856680255214,
635567,0.002334490699213771,0.8186117238835449]	97.0	(113, [97], [1.0])	[0.5689094981277257, 0.004796495596283756,
18911745,0.0033146713992235695,0.8819427857731233]	97.0	(113, [97], [1.0])	[0.44836134868724614, 0.0056525555745213524
36809091615388,0.8497142611454308]	97.0	(113, [97], [1.0])	[0.49981558447821167, 0.0073583961048181165
368256869812,0.8846311177872647]	97.0	(113, [97], [1.0])	[0.462417102342522, 0.00425015719064818, 0.
87626764338403,0.8477323474233768]	97.0	(113, [97], [1.0])	[0.5261711159039357, 0.0056308162701958495,

Cosine similarity score to determine the similarity between two songs.

```

from pyspark.sql.functions import udf
from pyspark.ml.linalg import Vectors, VectorUDT
from pyspark.sql.types import DoubleType
import numpy as np

# UDF to calculate the dot product of two vectors
def cosine_similarity(list_a, list_b):
    # Convert lists to numpy arrays
    np_a = np.array(list_a)
    np_b = np.array(list_b)
    # Calculate dot product using numpy
    #return float(np.dot(np_a, np_b)) #/ (np.linalg.norm(np_a) * np.linalg.norm(np_b))

    # Calculate dot product of features
    dot_product = np.dot(np_a, np_b)
    dot_product = float(dot_product)

    # Calculate the normalization of list a
    norm_vec1 = np.linalg.norm(list_a)
    norm_vec1 = float(norm_vec1)

    # Calculate the normalization of list b
    norm_vec2 = np.linalg.norm(list_b)
    norm_vec2 = float(norm_vec2)

    # Calculate the cosine similarity of all songs
    cosine_similarity = dot_product / (norm_vec1 * norm_vec2) if norm_vec1 != 0 and norm_vec2 != 0 else 0.0
    return cosine_similarity

cosine_similarity_udf = udf(cosine_similarity, DoubleType())

```

Recommend the top 10 songs based on a song track.

```

from pyspark.sql.functions import col, lit
import pyspark.sql.functions as F

# Assuming `song_id` is a unique identifier for each song
def recommend_songs(song_id, df):
    # Filter out the target song
    target_song = df.filter(df.track_id == song_id).select("features").collect()[0][0]
    #return target_song

    # Add a column for the target song's features
    df_with_similarity = df.withColumn("targetFeatures", F.array([F.lit(x) for x in target_song]))
    #df_with_similarity = df.withColumn("targetFeatures", lit(target_song))

    # Calculate cosine similarity between target song and all songs
    df_with_similarity = df_with_similarity.withColumn("similarity", cosine_similarity_udf("features", "targetFeatures"))

    # Sort by similarity and fetch top 10 songs
    df_filtered = df_with_similarity.filter(df_with_similarity.track_id != song_id)
    top_songs = df_filtered.sort(col("similarity").desc()).limit(10)

    return top_songs

# Example usage
song_id = '6DCZcSspjsKoFjzjrWoCdn'
top_recommendations = recommend_songs(song_id, df_transformed)
top_recommendations
top_recommendations.show(truncate=False)

```

```

-----
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
-----

```



```
df_transformed.filter(df_transformed.artists == 'Drake').show(truncate=False)
```

track_id	artists	album_name	track_name	popularity	explicit	danceability	energy	loudness	mode	speechiness	acou
6DCZcSspjsKoFjzjrWoCdn	Drake	Scorpion	God's Plan	84	true	0.754	0.449	-9.211	1	0.109	0.03

✓ Find similar songs

The function takes in a query and returns the data we want about a song to search for.

```
def search_songs_in_df(query):
    # Filter DataFrame where song_name contains the query string (case-insensitive)
    filtered_df = df_transformed.filter(df_transformed.track_name.contains(query)).select("track_id", "track_name", 'artists', '
    return filtered_df
```

```
import ipywidgets as widgets
from IPython.display import display, clear_output
```

```
# Input text and button for the search
text = widgets.Text(value='', placeholder='Type a song name', description='Search:', disabled=False)
button = widgets.Button(description="Search")
```

```
# Dropdown for selecting a song from search results
dropdown = widgets.Dropdown(options=['Select a song'], description='Results:', disabled=False)
```

```
# Output widget for displaying the track ID
output = widgets.Output()
```

```
display(text, button, dropdown, output)
```

```
def on_button_clicked(b):
    with output:
        clear_output()
        query = text.value
        if query:
            results_df = search_songs_in_df(query).collect()
            # Update dropdown options with search results
            dropdown.options = [(row.track_name, row.artists, row.album_name, row.track_id) for row in results_df]
            dropdown.disabled = False if results_df else True
        else:
            print("Please enter a query.")
```

```
def on_dropdown_change(change):
    with output:
        if change['name'] == 'value' and change['new']:
            clear_output()
            print(f"Track Name: {change['new'][0]}")
            print(f"Artist: {change['new'][1]}")
            print(f"Album: {change['new'][2]}")
            print(f"Track ID: {change['new'][3]}")

            song_id = change['new'][3]
            top_recommendations = recommend_songs(song_id, df_transformed)
            top_recommendations.select("track_name", "artists", "album_name").show(truncate=False)
```